

Final Project - Group 2

Sam, Eva, Dries, Caden, Jonas

2024-12-12

```
suppressMessages({  
  library(ggplot2)  
  library(reshape2)  
  library(dplyr)  
  library(faraway)  
  library(lubridate)  
})
```

```
## Warning in check_dep_version(): ABI version mismatch:  
## lme4 was built with Matrix ABI version 1  
## Current Matrix ABI version is 0  
## Please re-install lme4 from source or restore original 'Matrix' package  
## Warning: package 'lubridate' was built under R version 4.3.3
```

1 - Exploratory Data Analysis

1.11 - Testing WHOOP's Activity Strain Metric - LINEAR REGRESSION

```
workouts <- read.csv("data/workouts.csv")
```

```
workouts_1 <- workouts %>%  
  select(-`GPS.enabled`,  
    -`Distance..meters.`,  
    -`Altitude.gain..meters.`,  
    -`Altitude.change..meters.`,  
    -`Cycle.start.time`,  
    -`Cycle.end.time`,  
    -`Cycle.timezone`,  
    -`Workout.start.time`,  
    -`Workout.end.time`,  
    -`Activity.name`)
```

```
activity_strain <- lm(Activity.Strain ~ ., data = workouts_1)  
summary(activity_strain)
```

##	Estimate	Std. Error	t value	Pr(> t)
## (Intercept)	-1.95728417	0.75131850	-2.6051	0.009298
## Duration..min.	-0.00303960	0.00127355	-2.3867	0.017154
## Energy.burned..cal.	0.01237841	0.00023213	53.3260	< 2.2e-16
## Max.HR..bpm.	0.01801828	0.00226009	7.9724	3.603e-15
## Average.HR..bpm.	0.04057875	0.00787245	5.1545	2.972e-07

```
## HR.Zone.1..      -0.00065590  0.00307877 -0.2130  0.831332
## HR.Zone.2..      -0.01182269  0.00398742 -2.9650  0.003087
## HR.Zone.3..      -0.00244303  0.00560155 -0.4361  0.662817
## HR.Zone.4..       0.01077301  0.00720889  1.4944  0.135333
## HR.Zone.5..       0.02787134  0.01163991  2.3945  0.016798
```

```
##
```

```
## n = 1209, p = 10, Residual SE = 0.69567, R-Squared = 0.95
```

```
summary(activity_strain)$adj.r.squared
```


```
## [1] 0.9541699
```

Conclusion: 95% of the variability in Activity Strain is explained by our reduced dataset and a very simple model. This may gives us reason to drop this variable as well.

1.12 - Testing WHOOP's Activity Strain Metric - CORRELATION PLOT

```
cor_matrix <- cor(workouts_1, use = "complete.obs")
melted_cor_matrix <- melt(cor_matrix)

ggplot(melted_cor_matrix, aes(Var1, Var2, fill = value)) +
  geom_tile() +
  scale_fill_gradientn(colors = c("blue", "white", "red"), values = c(0, 0.5, 1), limit = c(-1, 1)) +
  theme_light(base_size = 14) +
  theme(axis.text.x = element_text(angle = 45, hjust = 1),
        axis.text.y = element_text(size = 12)) +
  labs(title = "Correlation Matrix Heatmap", x = "Variables", y = "Variables") +
  theme(plot.title = element_text(hjust = 0.5, size = 16, face = "bold"))
```



```
threshold <- 0.85
high_corr_pairs <- which(abs(cor_matrix) > threshold, arr.ind = TRUE)
high_corr_df <- data.frame(
  Column1 = rownames(cor_matrix)[high_corr_pairs[, 1]],
  Column2 = colnames(cor_matrix)[high_corr_pairs[, 2]],
  Correlation = cor_matrix[high_corr_pairs]
)
high_corr_df <- high_corr_df[high_corr_df$Column1 != high_corr_df$Column2, ]
print(high_corr_df)
```

```
##           Column1           Column2 Correlation
## 3 Energy.burned..cal. Activity.Strain 0.9256782
## 4 Activity.Strain Energy.burned..cal. 0.9256782
```

Conclusion: to avoid collinearity between Energy Burned and Activity Strain, we will drop the Activity Strain column since Energy Burned provides an empirical datapoint.

Conclusion: to avoid collinearity between Energy Burned and Activity Strain, we will drop the Activity Strain column since Energy Burned provides an empirical datapoint.

2 - Data Prep, Preparing the Explanatory Variables (Workouts)

2.1 - daily_workouts

```
##note: premise of time series = One Day One Observation

# Step 1: Drop the columns with no data
daily_workouts <- workouts %>%
  select(-c("GPS.enabled", "Distance..meters.",
            "Altitude.gain..meters.", "Altitude.change..meters."))

# Step 2: Remove the useless or redundant data
daily_workouts <- daily_workouts %>%
  select(-c("Cycle.start.time", "Cycle.end.time", "Cycle.timezone"))

# Step 3: Convert Activity.name to a factor
# Note: There was a level in the factor called "Other" with 1 observation,
# which was removed for clarity, since "Other" is no interpretable.
# Note: We also removed the activity "Activity" also for lack of interpretability.
# This activity had 65 observations.
daily_workouts$Activity.name <- as.factor(daily_workouts$Activity.name)
daily_workouts <- daily_workouts %>%
  mutate(
    Assault_Bike = Activity.name == "Assault Bike",
    Cycling = Activity.name == "Cycling",
    Elliptical = Activity.name == "Elliptical",
    Functional_Fitness = Activity.name == "Functional Fitness",
    Hiking_Rucking = Activity.name == "Hiking/Rucking",
    Powerlifting = Activity.name == "Powerlifting",
    Rowing = Activity.name == "Rowing",
```

```

Running = Activity.name == "Running",
Spin = Activity.name == "Spin",
Stairmaster = Activity.name == "Stairmaster",
Walking = Activity.name == "Walking",
Weightlifting = Activity.name == "Weightlifting",
Yoga = Activity.name == "Yoga"
) %>%
select(-Activity.name) # Remove the original Activity.name column

# Step 4: Convert 'Workout.start.time' to Date and extract the date portion
daily_workouts <- daily_workouts %>%
mutate(Workout.start.date = as.Date(Workout.start.time, format = "%Y-%m-%d %H:%M:%S"))

# Step 5: Group by date and summarize
daily_workouts <- daily_workouts %>%
group_by(Workout.start.date) %>%
summarise(
  # Sum of 'Duration..min.' and 'Energy.burned..cal.'
  Total.Duration.min = sum(`Duration..min.`, na.rm = TRUE),
  Total.Energy.burned.cal = sum(`Energy.burned..cal.`, na.rm = TRUE),

  # Weighted average time spent at average heart rate
  Weighted.Avg.HR = sum(`Average.HR..bpm.` * `Duration..min.`, na.rm = TRUE) /
    sum(`Duration..min.`, na.rm = TRUE),

  # Weighted average time spent for heart rate zones
  Weighted.Avg.HR.Zone.1 = sum(`HR.Zone.1..` * `Duration..min.`, na.rm = TRUE) /
    sum(`Duration..min.`, na.rm = TRUE),
  Weighted.Avg.HR.Zone.2 = sum(`HR.Zone.2..` * `Duration..min.`, na.rm = TRUE) /
    sum(`Duration..min.`, na.rm = TRUE),
  Weighted.Avg.HR.Zone.3 = sum(`HR.Zone.3..` * `Duration..min.`, na.rm = TRUE) /
    sum(`Duration..min.`, na.rm = TRUE),
  Weighted.Avg.HR.Zone.4 = sum(`HR.Zone.4..` * `Duration..min.`, na.rm = TRUE) /
    sum(`Duration..min.`, na.rm = TRUE),
  Weighted.Avg.HR.Zone.5 = sum(`HR.Zone.5..` * `Duration..min.`, na.rm = TRUE) /
    sum(`Duration..min.`, na.rm = TRUE),

  # Keep the Max HR for the day
  Max.HR.bpm = max(`Max.HR..bpm.`, na.rm = TRUE),

  # Sum of boolean values for each activity (to count the occurrences per day)
  Total.Assault_Bike = sum(Assault_Bike, na.rm = TRUE),
  Total.Cycling = sum(Cycling, na.rm = TRUE),
  Total.Elliptical = sum(Elliptical, na.rm = TRUE),
  Total.Functional_Fitness = sum(Functional_Fitness, na.rm = TRUE),
  Total.Hiking_Rucking = sum(Hiking_Rucking, na.rm = TRUE),
  Total.Powerlifting = sum(Powerlifting, na.rm = TRUE),
  Total.Rowing = sum(Rowing, na.rm = TRUE),
  Total.Running = sum(Running, na.rm = TRUE),
  Total.Spin = sum(Spin, na.rm = TRUE),
  Total.Stairmaster = sum(Stairmaster, na.rm = TRUE),
  Total.Walking = sum(Walking, na.rm = TRUE),
  Total.Weightlifting = sum(Weightlifting, na.rm = TRUE),

```

```

    Total.Yoga = sum(Yoga, na.rm = TRUE)) %>%

#Create a new column for heart zone 0
    mutate(Weighted.Avg.HR.Zone.0 = round((100 - Weighted.Avg.HR.Zone.1 - Weighted.Avg.HR.Zone.2 -
      Weighted.Avg.HR.Zone.3 - Weighted.Avg.HR.Zone.4 - Weighted.Avg.HR.Zone.5), 2))

```

2.2 - missing_dates_df

```

# Step 1: Identify the date range (min and max dates)
date_range <- seq.Date(min(daily_workouts$Workout.start.date),
  max(daily_workouts$Workout.start.date), by = "day")

# Step 2: Create a data frame for all dates in the range
missing_dates_df <- data.frame(
  Workout.start.date = date_range,
  Total.Duration.min = 0,
  Total.Energy.burned.cal = 0,
  Weighted.Avg.HR = 57, # Set resting heart rate
  Weighted.Avg.HR.Zone.1 = 0,
  Weighted.Avg.HR.Zone.2 = 0,
  Weighted.Avg.HR.Zone.3 = 0,
  Weighted.Avg.HR.Zone.4 = 0,
  Weighted.Avg.HR.Zone.5 = 0,
  Max.HR.bpm = 57, # Set resting heart rate for Max HR
  Total.Assault_Bike = 0,
  Total.Cycling = 0,
  Total.Elliptical = 0,
  Total.Functional_Fitness = 0,
  Total.Hiking_Rucking = 0,
  Total.Powerlifting = 0,
  Total.Rowing = 0,
  Total.Running = 0,
  Total.Spin = 0,
  Total.Stairmaster = 0,
  Total.Walking = 0,
  Total.Weightlifting = 0,
  Total.Yoga = 0,
  Weighted.Avg.HR.Zone.0 = 100 # Set HR Zone 0 to 100
)

```

2.3 - Combining daily_workouts and missing_dates_df into One Continuous Exploratory Dataset

```

# Step 3: Combine with the original data
daily_workouts <- bind_rows(
  daily_workouts,
  missing_dates_df
)

# Step 4: Remove any duplicates and sort by date
daily_workouts <- daily_workouts %>%
  distinct(Workout.start.date, .keep_all = TRUE) %>%
  arrange(Workout.start.date)

```

Columns in wrangled workouts dataset:

- Workout.start.date
 - Total.Duration.min
 - Total.Energy.burned.cal
 - Weighted.Avg.HR
 - Weighted.Avg.HR.Zone.0
 - Weighted.Avg.HR.Zone.1
 - Weighted.Avg.HR.Zone.2
 - Weighted.Avg.HR.Zone.3
 - Weighted.Avg.HR.Zone.4
 - Weighted.Avg.HR.Zone.5
 - Max.HR.bpm
 - Total.Assault_Bike
 - Total.Cycling
 - Total.Elliptical
 - Total.Functional_Fitness
 - Total.Hiking_Rucking
 - Total.Powerlifting
 - Total.Rowing
 - Total.Running
 - Total.Spin
 - Total.Stairmaster
 - Total.Walking
 - Total.Weightlifting
 - Total.Yoga
-

3 - Data Prep, Preparing the Response Variable: Sleep

Columns in sleep dataset:

- Cycle.start.time
- Cycle.end.time
- Cycle.timezone
- Sleep.onset
- Wake.onset
- Sleep.performance..
- Respiratory.rate..rpm.
- Asleep.duration..min.
- In.bed.duration..min.
- Light.sleep.duration..min.
- Deep..SWS..duration..min.
- REM.duration..min.
- Awake.duration..min.
- Sleep.need..min.
- Sleep.debt..min.
- Sleep.efficiency..
- Sleep.consistency..
- Nap

3.1 - Sleep CORRELATION PLOT

```
# Load necessary libraries
library(ggplot2)
library(reshape2)


# Read the CSV file
sleeps <- read.csv("data/sleeps.csv")

# Select only numeric columns
sleeps_numeric <- sleeps[sapply(sleeps, is.numeric)]

# Compute the correlation matrix using complete observations
cor_matrix2 <- cor(sleeps_numeric, use = "complete.obs")

# Melt the correlation matrix for visualization
melted_cor_matrix2 <- melt(cor_matrix2)

# Create the heatmap using ggplot2
ggplot(melted_cor_matrix2, aes(Var1, Var2, fill = value)) +
  geom_tile() +
  scale_fill_gradientn(colors = c("blue", "white", "red"), values = c(0, 0.5, 1), limit = c(-1, 1)) +
  theme_light(base_size = 14) +
  theme(axis.text.x = element_text(angle = 45, hjust = 1),
        axis.text.y = element_text(size = 12)) +
  labs(title = "Correlation Matrix Heatmap", x = "Variables", y = "Variables") +
  theme(plot.title = element_text(hjust = 0.5, size = 16, face = "bold"))
```



```
# Define the correlation threshold
threshold <- 0.85

# Find pairs of variables with correlations exceeding the threshold
high_corr_pairs2 <- which(abs(cor_matrix2) > threshold, arr.ind = TRUE)

# Create a data frame of the high-correlation pairs
high_corr_df2 <- data.frame(
  Column1 = rownames(cor_matrix2)[high_corr_pairs2[, 1]],
  Column2 = colnames(cor_matrix2)[high_corr_pairs2[, 2]],
  Correlation = cor_matrix2[high_corr_pairs2]
)

# Remove self-correlations (diagonal elements)
high_corr_df2 <- high_corr_df2[high_corr_df2$Column1 != high_corr_df2$Column2, ]

# Print the high-correlation pairs
print(high_corr_df2)
```

```
##           Column1           Column2 Correlation
## 2  Asleep.duration..min.  Sleep.performance..  0.9266651
## 3  In.bed.duration..min.  Sleep.performance..  0.8528991
## 5    Sleep.performance.. Asleep.duration..min.  0.9266651
## 7  In.bed.duration..min. Asleep.duration..min.  0.8900266
## 8    Sleep.performance.. In.bed.duration..min.  0.8528991
## 9  Asleep.duration..min. In.bed.duration..min.  0.8900266
## 15   Sleep.efficiency..  Awake.duration..min. -0.8693928
## 18  Awake.duration..min.  Sleep.efficiency.. -0.8693928
```



```

#Step 1: remove unnecessary or redundant variables
daily_sleep <- sleeps %>%
  select(- `Cycle.start.time`,
         - `Cycle.end.time`,
         - `Cycle.timezone`,
         - `Sleep.performance..`,
         - `In.bed.duration..min.`,
         - `Sleep.debt..min.`,
         - `Sleep.efficiency..`,
         - `Nap`)

# Step 2: Group by Wake.date and calculate the number of wake-ups and total sleep time
nap_data <- daily_sleep %>%
  mutate(Wake.date = as.Date(Wake.onset)) %>%
  group_by(Wake.date) %>%
  summarise(
    Nap.count = n_distinct(Wake.onset) - 1,
    Nap.duration = sum(Asleep.duration..min., na.rm = TRUE) - max(Asleep.duration..min.)
  ) %>%
  ungroup() %>%
  mutate(Nap.duration = ifelse(Nap.count == 0, 0, Nap.duration))

# Step 3: Group by Wake.onset date and select the row with the largest Asleep.duration..min.
sleep_data <- daily_sleep %>%
  # Group by Wake.onset
  group_by(as.Date(Wake.onset)) %>%
  # Filter to keep the row with the max Asleep.duration..min.
  filter(Asleep.duration..min. == max(Asleep.duration..min.)) %>%
  slice(1) %>% ungroup()

#Step 4: We will think of your sleep as the day that you worked out
sleep_data <- sleep_data %>%
  mutate(Sleep.date = as.Date(Wake.onset) - days(1))

# Step 5: Merge sleep_data and nap_data by their respective date columns
sleep_nap_data <- sleep_data %>%
  left_join(nap_data, by = c("Sleep.date" = "Wake.date")) %>%
  filter(!is.na(Nap.count))

```

4 - Final Dataset, Merge Sleep and Workout Data

```
merged_data <- sleep_nap_data %>%
  inner_join(daily_workouts, by = c("Sleep.date" = "Workout.start.date")) %>%
  select(-`Sleep.onset`,
         - `Wake.onset`,
         - `as.Date(Wake.onset)`)

str(merged_data)

## tibble [673 x 34] (S3: tbl_df/tbl/data.frame)
##  $ Respiratory.rate..rpm.      : num [1:673] 13.8 14.3 14.1 13.6 14.3 13.9 14.8 14.2 14.4 13.9 ...
##  $ Asleep.duration..min.       : int [1:673] 444 422 286 321 369 378 502 428 412 536 ...
##  $ Light.sleep.duration..min.: int [1:673] 172 175 75 81 78 121 153 125 188 234 ...
##  $ Deep..SWS..duration..min. : int [1:673] 93 52 81 102 83 91 101 88 93 90 ...
##  $ REM.duration..min.         : int [1:673] 179 195 130 138 208 166 248 215 131 212 ...
##  $ Awake.duration..min.       : int [1:673] 44 78 55 44 93 150 126 31 52 34 ...
##  $ Sleep.need..min.           : int [1:673] 546 545 512 606 611 609 542 488 450 474 ...
##  $ Sleep.consistency..        : int [1:673] NA NA 67 78 66 74 58 60 61 46 ...
##  $ Sleep.date                 : Date[1:673], format: "2022-12-11" "2022-12-12" ...
##  $ Nap.count                  : num [1:673] 0 0 1 0 0 0 1 0 1 1 ...
##  $ Nap.duration               : num [1:673] 0 0 37 0 0 0 58 0 84 38 ...
##  $ Total.Duration.min         : num [1:673] 78 113 41 130 82 243 131 0 206 281 ...
##  $ Total.Energy.burned.cal     : num [1:673] 416 483 376 722 621 ...
##  $ Weighted.Avg.HR            : num [1:673] 129 119 155 129 145 ...
##  $ Weighted.Avg.HR.Zone.1     : num [1:673] 26.13 18.48 0 22.86 0.22 ...
##  $ Weighted.Avg.HR.Zone.2     : num [1:673] 28.4 53.3 2 34.8 26.8 ...
##  $ Weighted.Avg.HR.Zone.3     : num [1:673] 37.8 18.3 29 35.1 53.6 ...
##  $ Weighted.Avg.HR.Zone.4     : num [1:673] 0.769 0 66 4.377 19.463 ...
##  $ Weighted.Avg.HR.Zone.5     : num [1:673] 0 0 0 0 0 0 0 0 0 0 ...
##  $ Max.HR.bpm                 : num [1:673] 166 162 171 181 183 177 176 57 158 163 ...
##  $ Total.Assault_Bike         : num [1:673] 1 0 0 0 0 0 0 0 0 1 ...
##  $ Total.Cycling              : num [1:673] 0 0 0 0 0 0 0 0 0 0 ...
##  $ Total.Elliptical           : num [1:673] 0 0 0 0 0 0 0 0 0 0 ...
##  $ Total.Functional_Fitness   : num [1:673] 0 0 0 0 0 0 0 0 0 0 ...
##  $ Total.Hiking_Rucking       : num [1:673] 0 0 0 0 0 0 0 0 0 0 ...
##  $ Total.Powerlifting          : num [1:673] 0 0 0 0 0 0 0 0 0 0 ...
##  $ Total.Rowing               : num [1:673] 0 2 0 3 2 2 2 0 1 3 ...
##  $ Total.Running              : num [1:673] 1 0 1 1 0 0 0 0 0 0 ...
##  $ Total.Spin                 : num [1:673] 0 0 0 0 0 0 0 0 0 0 ...
##  $ Total.Stairmaster          : num [1:673] 0 0 0 0 0 0 0 0 0 0 ...
##  $ Total.Walking              : num [1:673] 0 0 0 0 0 0 0 0 0 0 ...
##  $ Total.Weightlifting        : num [1:673] 1 0 0 0 0 0 0 0 1 0 ...
##  $ Total.Yoga                 : num [1:673] 0 0 0 0 0 0 0 0 0 0 ...
##  $ Weighted.Avg.HR.Zone.0     : num [1:673] 6.9 9.93 3 2.84 0 ...

# Load necessary libraries
library(dplyr)

# Select the specified variables from merged_data_1
merged_data_1 <- merged_data %>%
  select(
    -Sleep.date,
    -Total.Assault_Bike,
```

```

-Total.Cycling,
-Total.Elliptical,
-Total.Functional_Fitness,
-Total.Hiking_Rucking,
-Total.Powerlifting,
-Total.Rowing,
-Total.Running,
-Total.Spin,
-Total.Stairmaster,
-Total.Walking,
-Total.Weightlifting,
-Total.Yoga
)

```

5 - Baseline Model

```

# Remove rows with missing values
cleaned_data <- na.omit(merged_data_1)

# Select only numeric columns
cleaned_numeric <- cleaned_data[sapply(cleaned_data, is.numeric)]


# Compute the correlation matrix using complete observations
cor_matrix3 <- cor(cleaned_numeric, use = "complete.obs")

# Melt the correlation matrix for visualization
melted_cor_matrix3 <- melt(cor_matrix3)

ggplot(melted_cor_matrix3, aes(Var1, Var2, fill = value)) +
  geom_tile() +
  scale_fill_gradientn(colors = c("blue", "white", "red"), values = c(0, 0.5, 1), limit = c(-1, 1)) +
  theme_light(base_size = 14) +
  theme(axis.text.x = element_text(angle = 45, hjust = 1),
        axis.text.y = element_text(size = 12)) +
  labs(title = "Correlation Matrix Heatmap", x = "Variables", y = "Variables") +
  theme(plot.title = element_text(hjust = 0.5, size = 16, face = "bold"))

```

Correlation Matrix Heatmap



```
# Define the correlation threshold
threshold <- 0.85

high_corr_pairs3 <- which(abs(cor_matrix3) > threshold, arr.ind = TRUE)
high_corr_df3 <- data.frame(
  Column1 = rownames(cor_matrix3)[high_corr_pairs3[, 1]],
  Column2 = colnames(cor_matrix3)[high_corr_pairs3[, 2]],
  Correlation = cor_matrix3[high_corr_pairs3]
)
high_corr_df3 <- high_corr_df3[high_corr_df3$Column1 != high_corr_df3$Column2, ]


print(high_corr_df3)
```

##	Column1	Column2	Correlation
## 12	Total.Energy.burned.cal	Total.Duration.min	0.8937230
## 13	Total.Duration.min	Total.Energy.burned.cal	0.8937230
## 16	Max.HR.bpm	Weighted.Avg.HR	0.9418197
## 17	Weighted.Avg.HR.Zone.0	Weighted.Avg.HR	-0.9730197
## 23	Weighted.Avg.HR	Max.HR.bpm	0.9418197
## 25	Weighted.Avg.HR.Zone.0	Max.HR.bpm	-0.9413541
## 26	Weighted.Avg.HR	Weighted.Avg.HR.Zone.0	-0.9730197
## 27	Max.HR.bpm	Weighted.Avg.HR.Zone.0	-0.9413541


```
# Basic model, no interaction terms
model <- lm(Deep..SWS..duration..min./Asleep.duration..min. ~ ., data = cleaned_data)
summary(model)$adj.r.squared

## [1] 0.3470618
```

```
plot(model, 1)
```



```
plot(model, 2)
```




```
# Model with Interactive terms (note this actually has a stronger R^2 than the highly complex model)
model_rr <- lm(Deep..SWS..duration..min./Asleep.duration..min. ~ (Total.Energy.burned.cal)*(.), data = )
```


```
summary(model_rr)$adj.r.squared
```

```
## [1] 0.3509511
```

```
plot(model_rr, 1)
```



```
plot(model_rr, 2)
```




lm(Deep..SWS..duration..min./Asleep.duration..min. ~ (Total.Energy.burned.c ...


```
# highly complex model
model_full <- lm(Deep..SWS..duration..min./Asleep.duration..min. ~ (.)^2, data = cleaned_data)
summary(model_full)$adj.r.squared

## [1] 0.3643356

plot(model_full, 1)
```



```
plot(model_full, 2)
```



6 - Model Selection

#Originally the model_full was not running on posit - we examined the dataset again and were able to ta

Results with Interactive term were:

Start: AIC=-4415.95)

Results with full model term were:

Start: AIC=-4295.39)

Remove rows with missing values

`cleaned_data <- na.omit(merged_data_1)`

`model_backward <- step(model_rr, direction = "backward")`

Intercept-only model with the cleaned data

`model_intercept <- lm(Deep..SWS..duration..min./Asleep.duration..min. ~ 1, data = cleaned_data)`

`model_forward <- step(model_intercept, direction = "forward", scope = formula(model_rr))`

#Best Adj-R²

`formula(model_backward)`

Deep..SWS..duration..min./Asleep.duration..min. ~ Total.Energy.burned.cal +

Light.sleep.duration..min. + REM.duration..min. + Awake.duration..min. +

Sleep.need..min. + Sleep.consistency.. + Total.Duration.min +

Weighted.Avg.HR + Weighted.Avg.HR.Zone.1 + Weighted.Avg.HR.Zone.2 +

Weighted.Avg.HR.Zone.3 + Max.HR.bpm + Weighted.Avg.HR.Zone.0 +

Total.Energy.burned.cal:Total.Duration.min + Total.Energy.burned.cal:Weighted.Avg.HR +

Total.Energy.burned.cal:Weighted.Avg.HR.Zone.1 + Total.Energy.burned.cal:Weighted.Avg.HR.Zone.2 +


Total.Energy.burned.cal:Weighted.Avg.HR.Zone.3 + Total.Energy.burned.cal:Max.HR.bpm +

Total.Energy.burned.cal:Weighted.Avg.HR.Zone.0

`summary(model_backward)$adj.r.squared`


[1] 0.3626397

`plot(model_backward, 1)`



lm(Deep..SWS..duration..min./Asleep.duration..min. ~ Total.Energy.burned.ca ...

```
plot(model_backward, 2)
```



lm(Deep..SWS..duration..min./Asleep.duration..min. ~ Total.Energy.burned.ca ...

```
formula(model_forward)
```


```
## Deep..SWS..duration..min./Asleep.duration..min. ~ Light.sleep.duration..min. +
## Max.HR.bpm + Sleep.consistency.. + Weighted.Avg.HR.Zone.4 +
```

```
## Awake.duration..min. + REM.duration..min.
```


```
summary(model_forward)$adj.r.squared
```

```
## [1] 0.3525724
```

```
plot(model_forward, 1)
```



```
plot(model_forward, 2)
```



7 - Assumption Checking for Linear Models

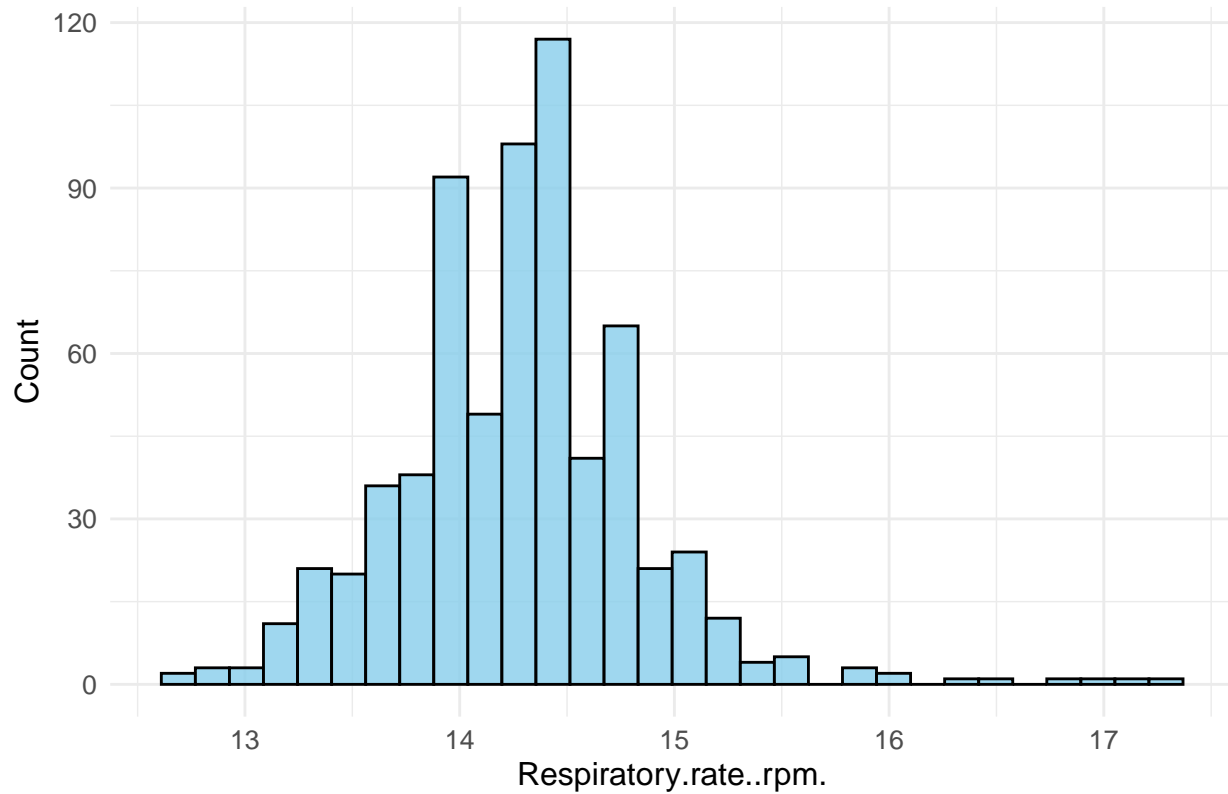
```
library(ggplot2)

predictors <- merged_data_1 %>%
  select(-Deep..SWS..duration..min.)


for (predictor in names(predictors)) {
  print(
    ggplot(data = predictors, aes_string(x = predictor)) +
      geom_histogram(fill = "skyblue", color = "black", bins = 30, alpha = 0.8) +
      theme_minimal() +
      labs(title = paste("Distribution of", predictor),
           x = predictor, y = "Count") +
      theme(
        plot.title = element_text(size = 14, face = "bold"),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10)
      )
  )
}

## Warning: `aes_string()` was deprecated in ggplot2 3.0.0.
## i Please use tidy evaluation idioms with `aes()`.
## i See also `vignette("ggplot2-in-packages")` for more information.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.
```

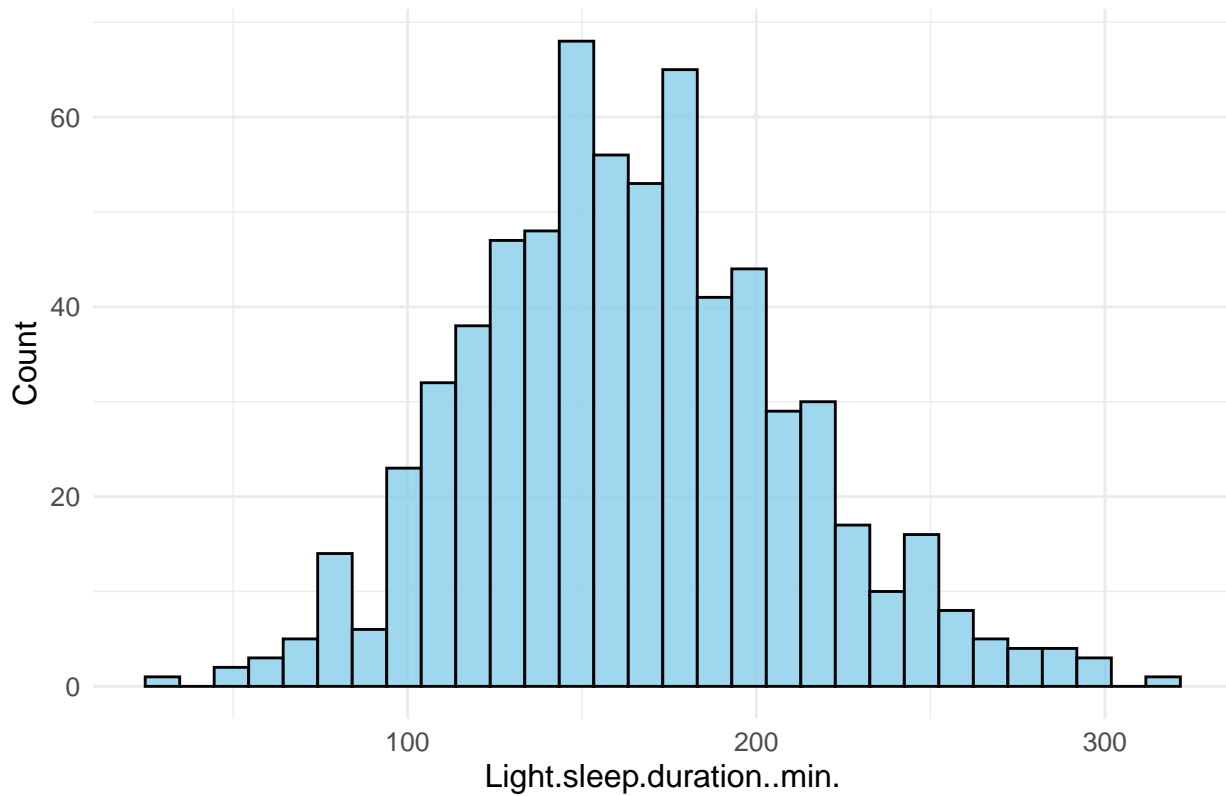
Distribution of Respiratory.rate..rpm.



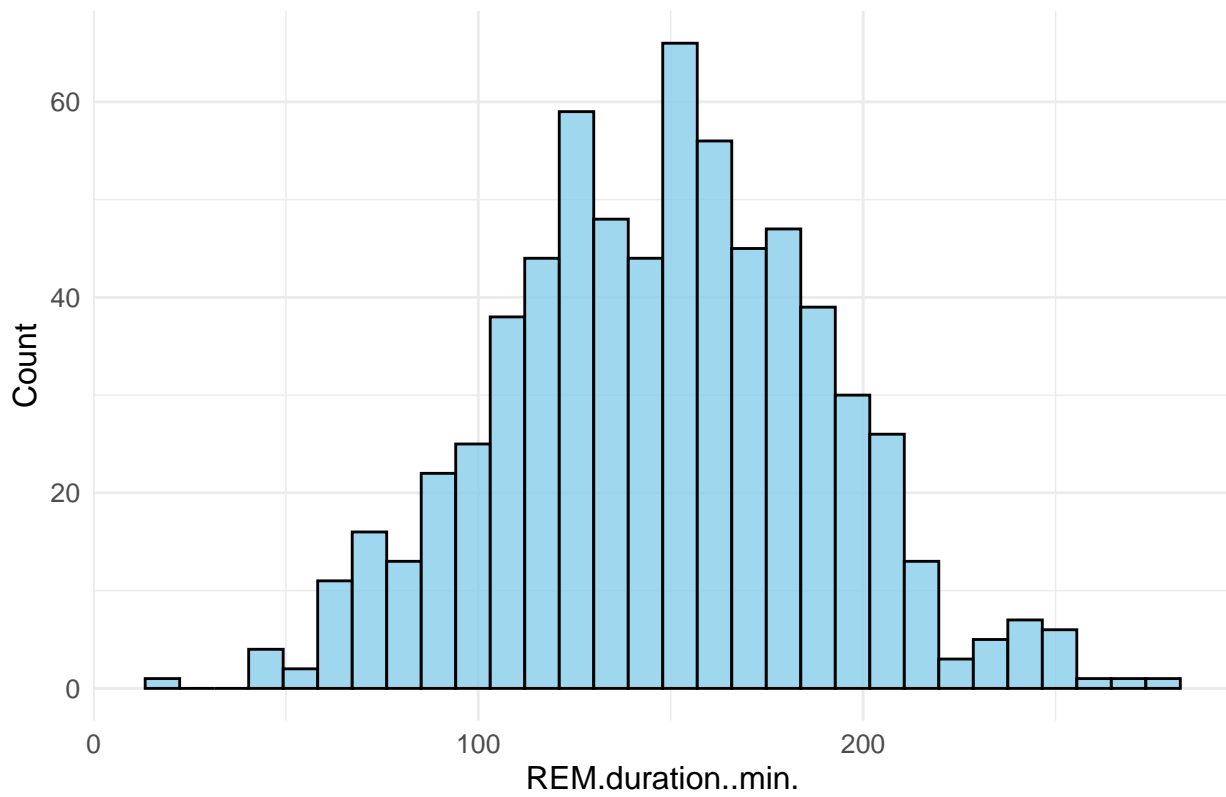
Distribution of Asleep.duration..min.




Distribution of Light.sleep.duration..min.



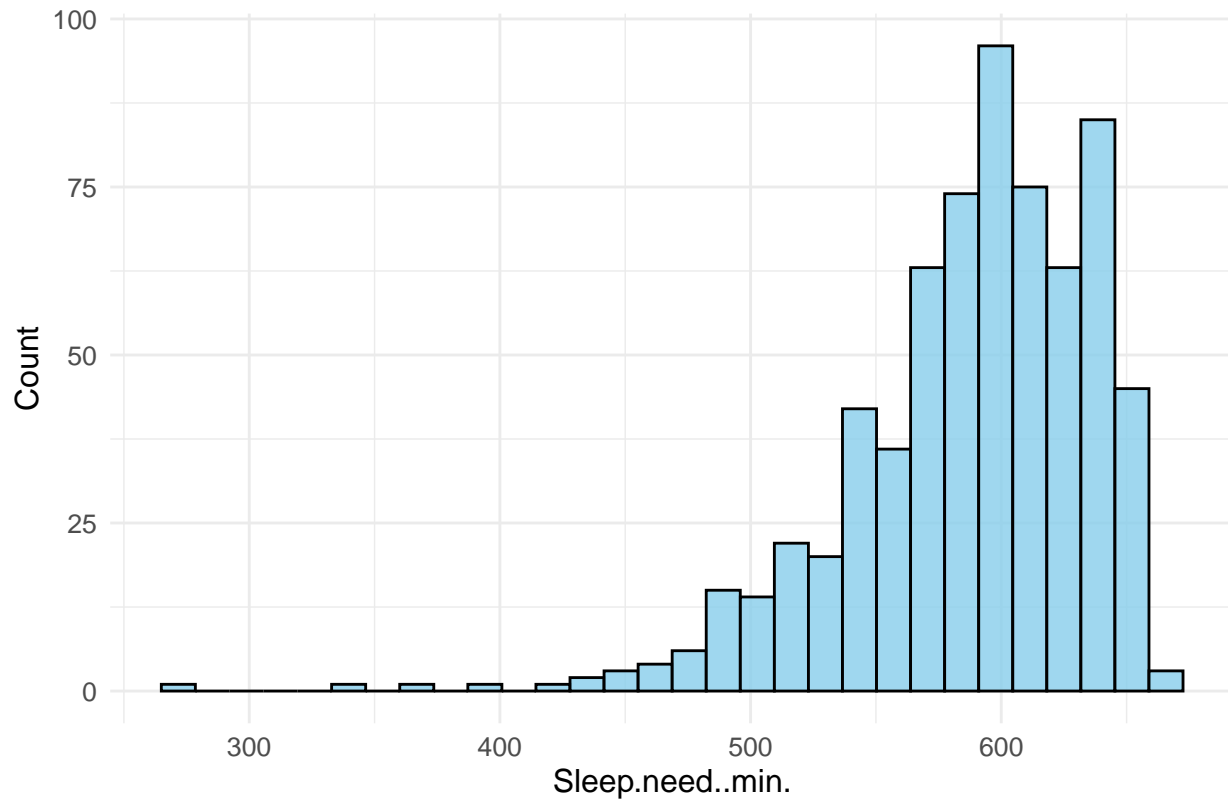
Distribution of REM.duration..min.



Distribution of Awake.duration..min.



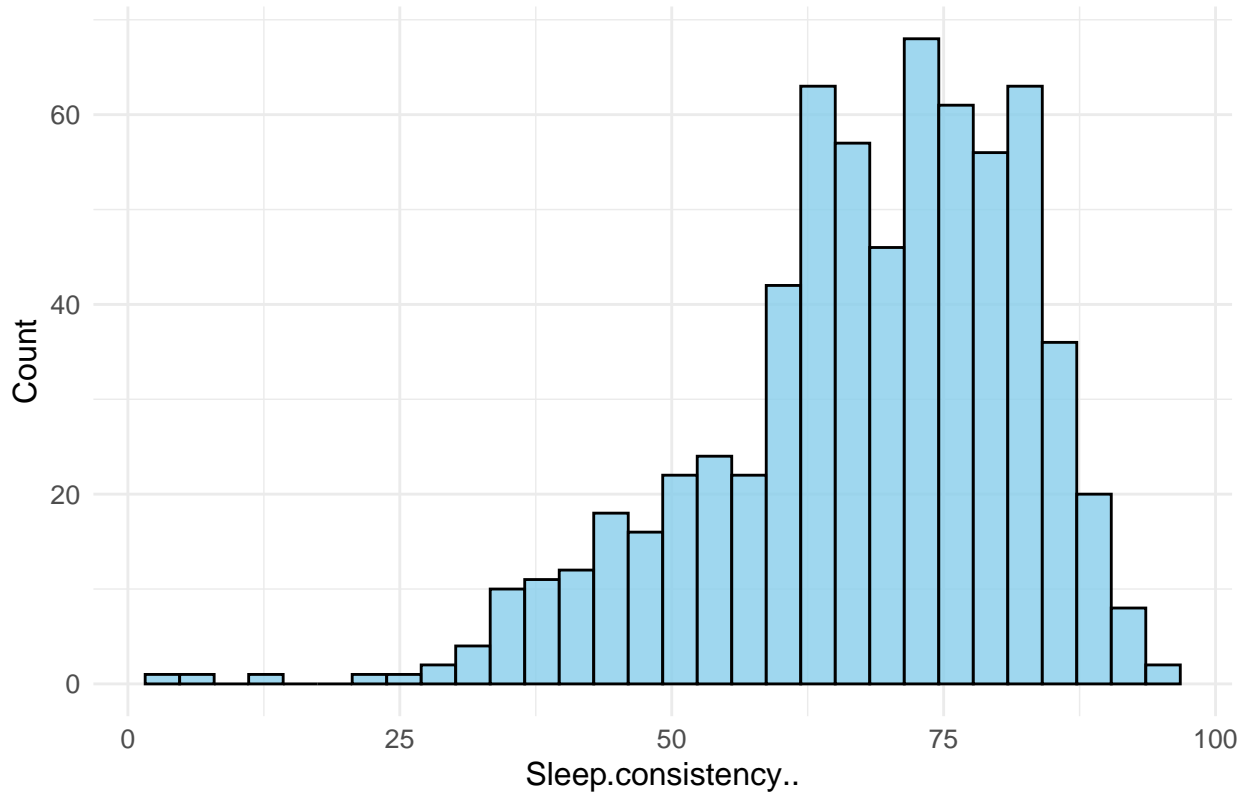
Distribution of Sleep.need..min.



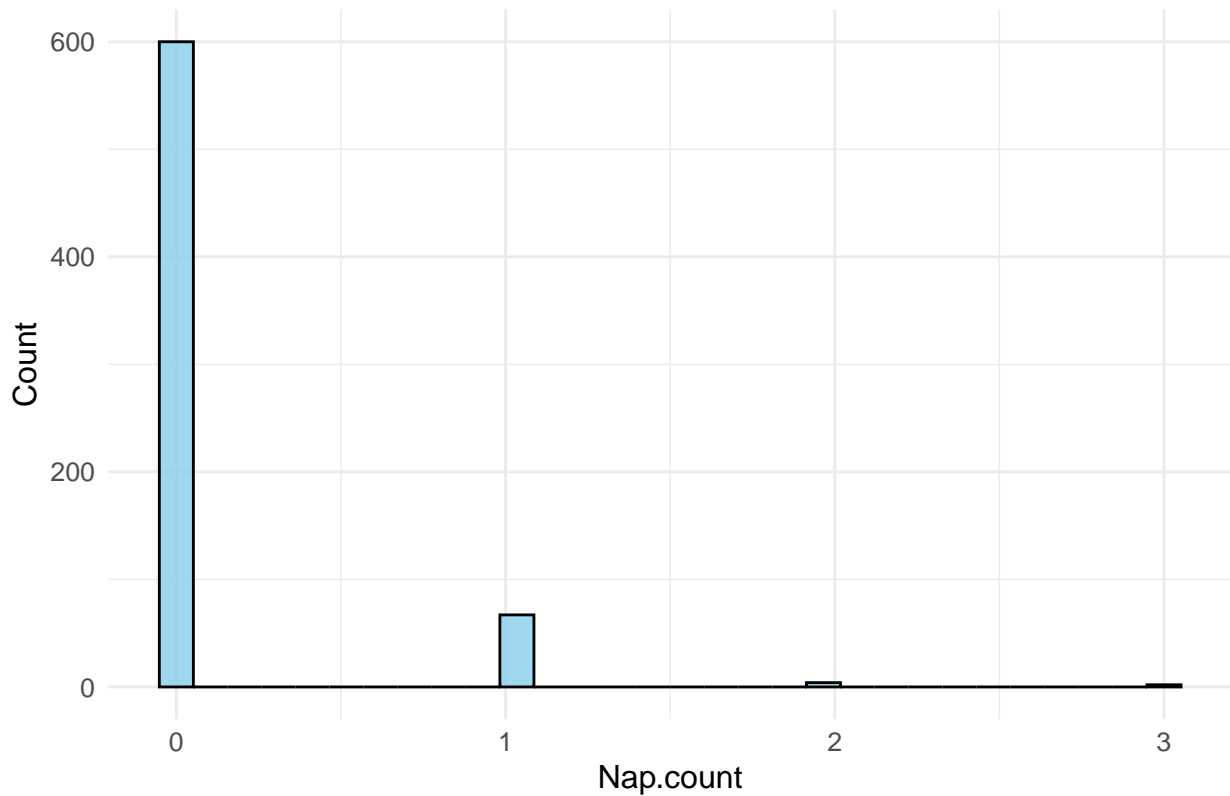
Warning: Removed 5 rows containing non-finite outside the scale range

```
## (`stat_bin()`).
```


Distribution of Sleep.consistency..



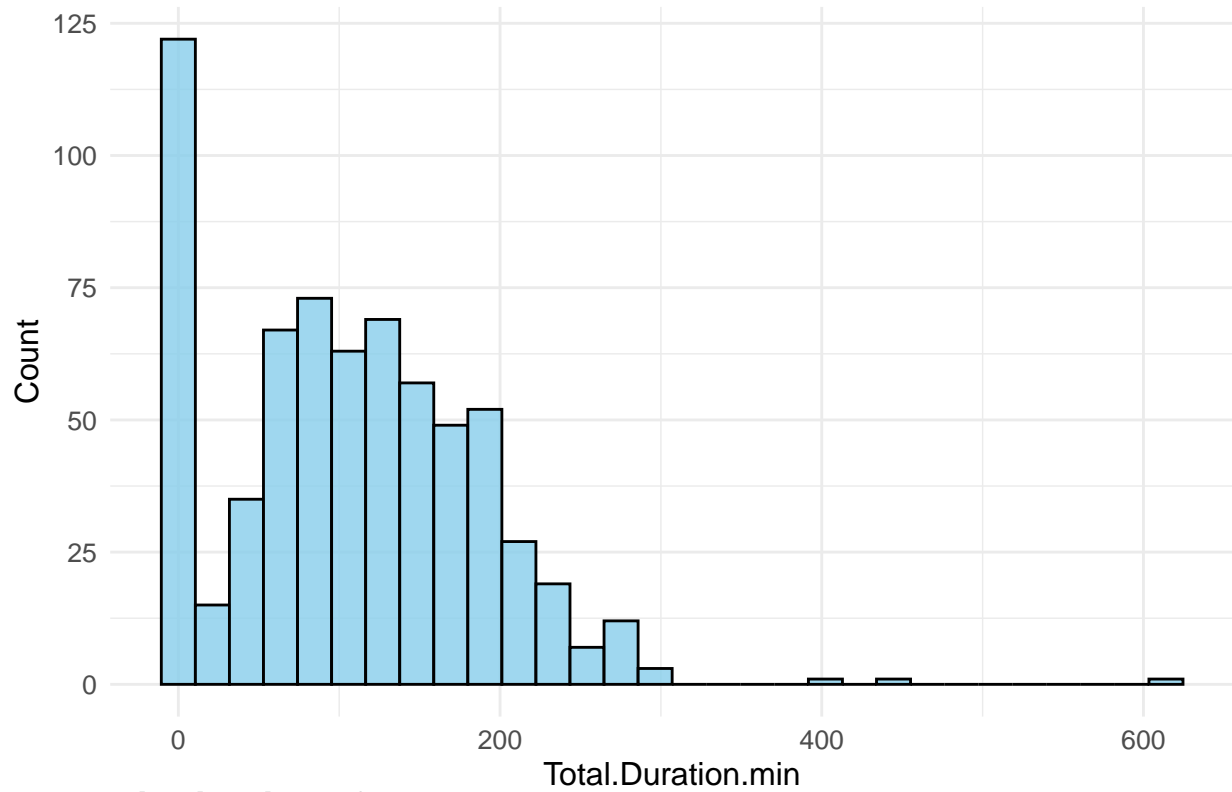
Distribution of Nap.count



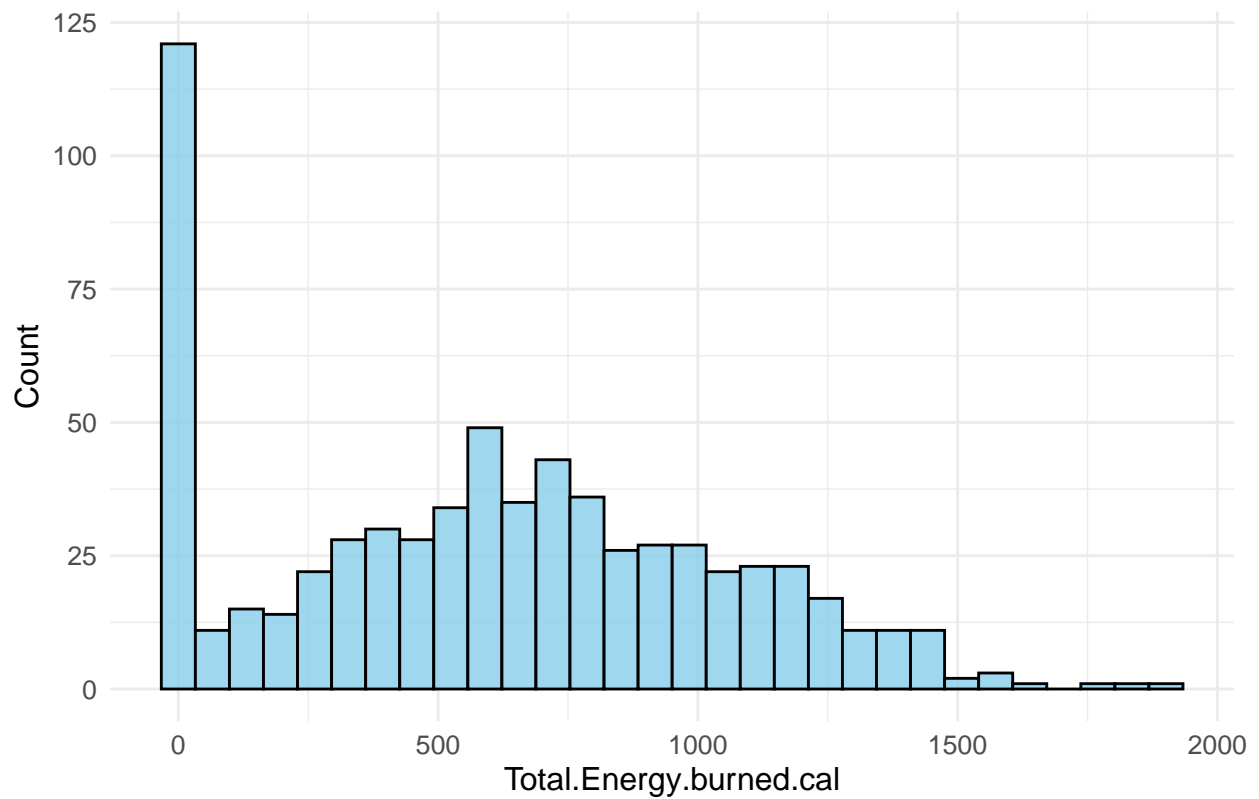

```
## Warning: Removed 2 rows containing non-finite outside the scale range
## (`stat_bin()`).
```



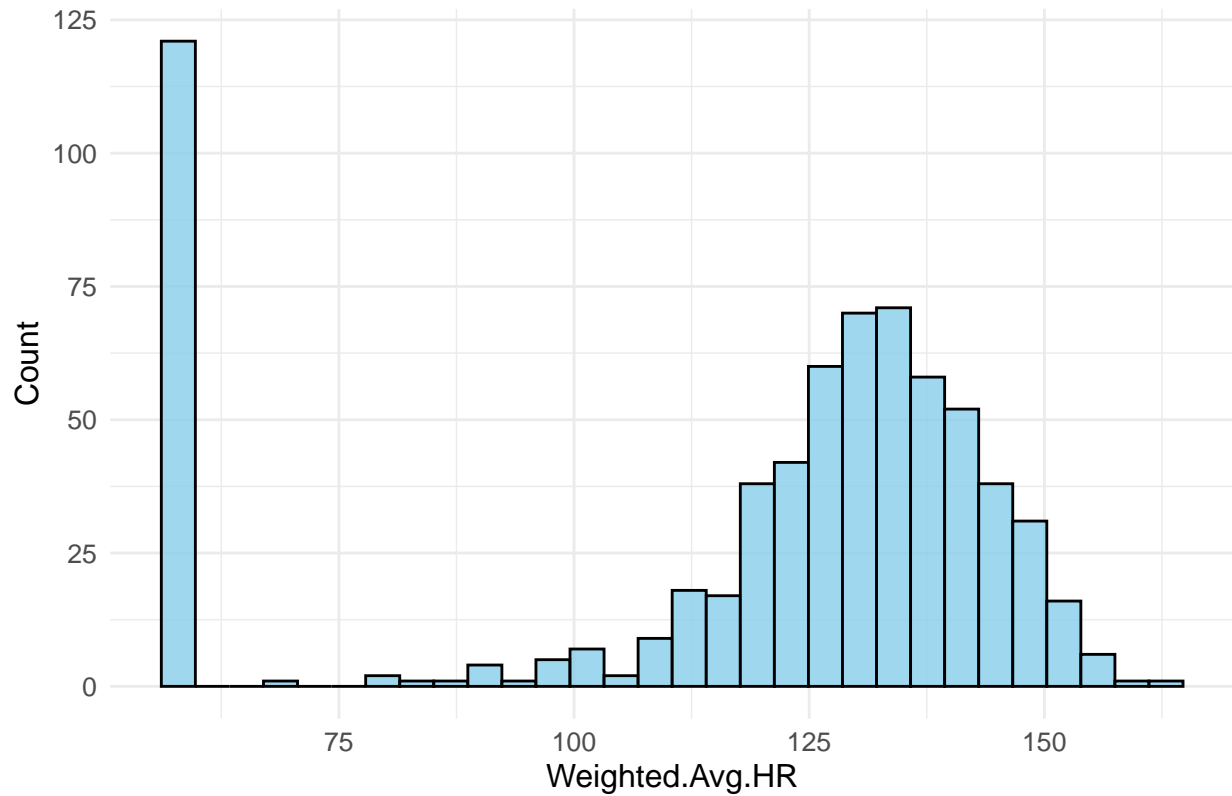
Distribution of Total.Duration.min




Distribution of Total.Energy.burned.cal



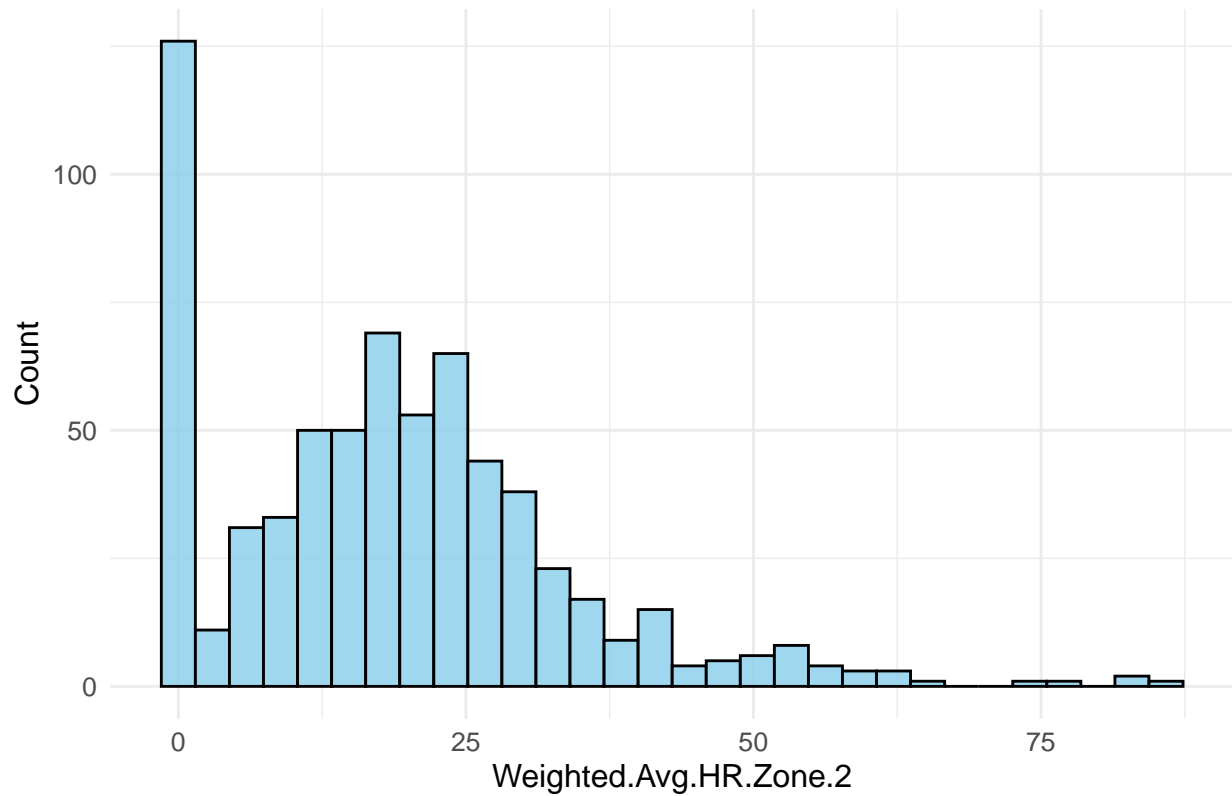
Distribution of Weighted.Avg.HR




Distribution of Weighted.Avg.HR.Zone.1



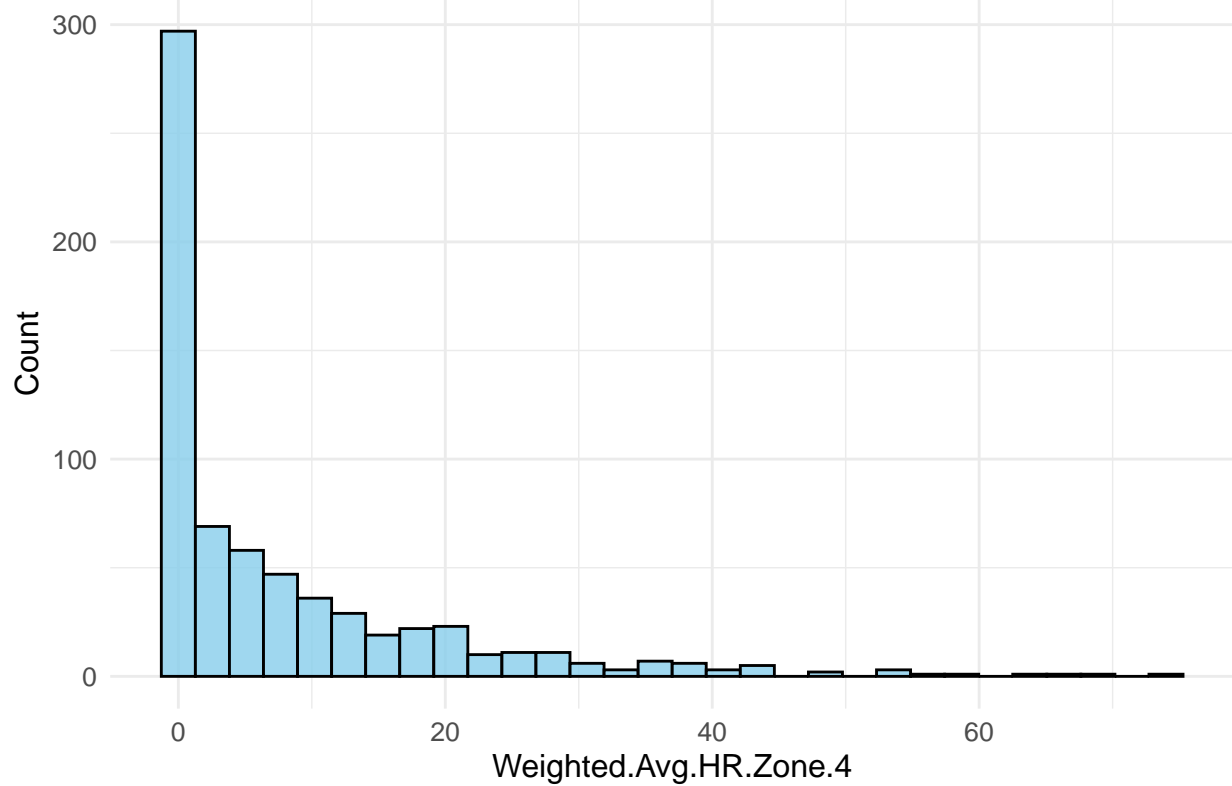
Distribution of Weighted.Avg.HR.Zone.2




Distribution of Weighted.Avg.HR.Zone.3



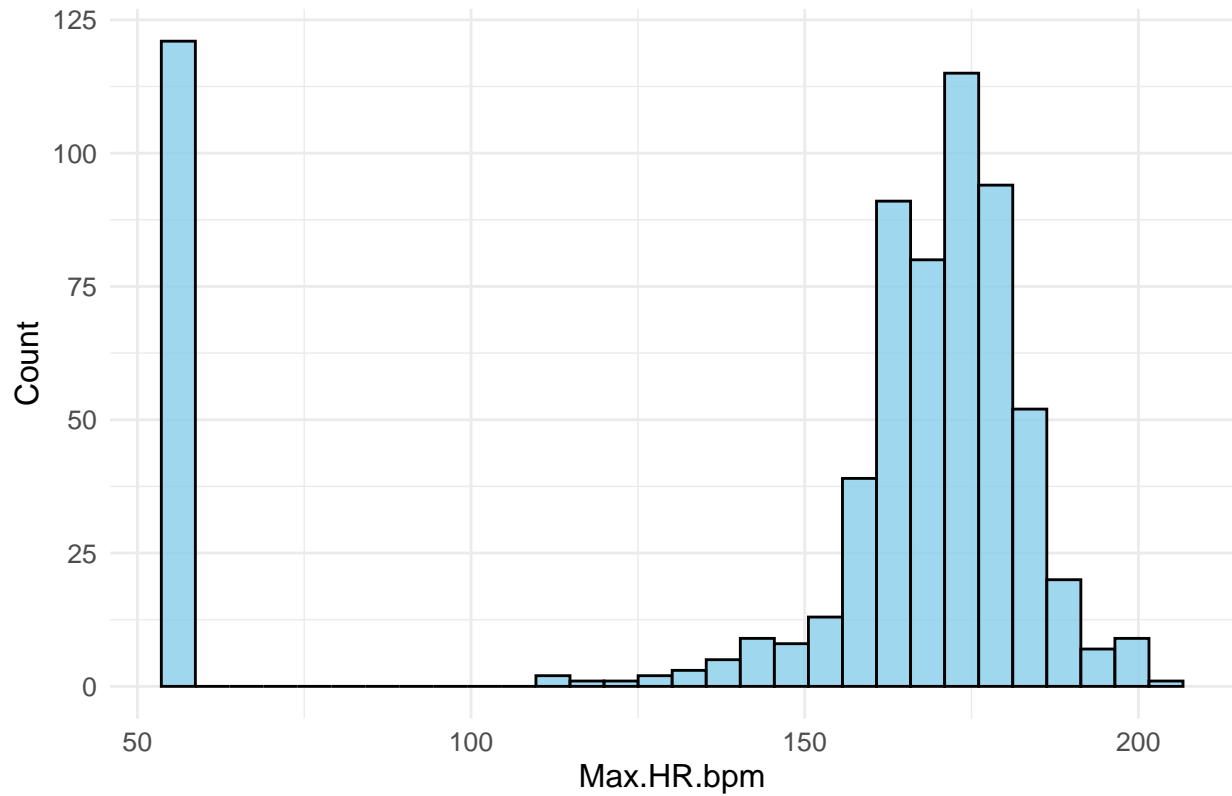
Distribution of Weighted.Avg.HR.Zone.4



Distribution of Weighted.Avg.HR.Zone.5




Distribution of Max.HR.bpm



Distribution of Weighted.Avg.HR.Zone.0



```
# 1. Linearity and Homoscedasticity - Residuals vs Fitted
par(mfrow = c(1, 2)) # Split screen into 2 plots
plot(model_backward, which = 1) # Residuals vs Fitted
plot(model_backward, which = 2) # Q-Q plot for residuals
```




```
# Extract the model's data for consistent lengths
model_data <- model.frame(model_backward)
residuals_backward <- residuals(model_backward)
```


```
# 2. Independence - Residuals over Total Duration
plot(model_data$Total.Duration.min, residuals_backward,
     main = "Residuals vs Total Duration",
     xlab = "Total Duration (min)", ylab = "Residuals",
     col = "darkblue", pch = 20)
```

```
# 3. Residual Histogram
hist(residuals_backward, breaks = 30, col = "skyblue",
     main = "Histogram of Residuals", xlab = "Residuals")
```

Residuals vs Total Duration



Histogram of Residuals



8 - Prediction

8.1 - Baseline

```
# Why? A simple model to establish a baseline for predictive performance.
# Linear regression helps identify if the relationship between predictors and response is linear. We can
# use a linear model to predict the response variable.

library("Metrics")
set.seed(139)
split <- sample(1:nrow(merged_data_1), 0.7 * nrow(merged_data_1)) # 70/30 split

train <- merged_data_1[split, ]
test <- merged_data_1[-split, ]

# Baseline Linear Regression
baseline_model <- lm(Deep..SWS..duration..min. / Asleep.duration..min. ~ ., data = train)
baseline_predictions <- predict(baseline_model, test)

# Model evaluation (RMSE and R-squared)
library(Metrics)
rmse_baseline <- rmse(test$Deep..SWS..duration..min. / test$Asleep.duration..min., baseline_predictions)
r2_baseline <- cor(test$Deep..SWS..duration..min. / test$Asleep.duration..min., baseline_predictions)^2

print(rmse_baseline)

## [1] NA
print(r2_baseline)
```



```
## [1] NA
```

8.2 - LASSO

```
# Lasso Regression  
# Why? Useful for datasets with many predictors to automatically perform feature selection.  
# Lasso helps reduce overfitting and simplifies the model by shrinking less important coefficients to zero  
  
set.seed(139)  
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-8
```

```
library(Metrics)  
  
train_clean <- na.omit(train)  
test_clean <- na.omit(test)  
  
x_train <- as.matrix(train_clean %>% select(-Deep..SWS..duration..min.))  
y_train <- train_clean$Deep..SWS..duration..min. / train_clean$Asleep.duration..min.  
  
x_test <- as.matrix(test_clean %>% select(-Deep..SWS..duration..min.))  
y_test <- test_clean$Deep..SWS..duration..min. / test_clean$Asleep.duration..min.  
  
# Fit Lasso model  
lasso_model <- cv.glmnet(x_train, y_train, alpha = 1)  
best_lambda <- lasso_model$lambda.min  
  
# Predict on test set  
lasso_preds <- predict(lasso_model, s = best_lambda, newx = x_test)  
  
# Calculate RMSE  
rmse_lasso <- rmse(y_test, lasso_preds)  
  
# Calculate R-squared  
rss <- sum((y_test - lasso_preds)^2)  
tss <- sum((y_test - mean(y_test))^2)  
r2_lasso <- 1 - (rss / tss)  
  
print(paste("Lasso RMSE:", round(rmse_lasso, 4)))
```

```
## [1] "Lasso RMSE: 0.0085"
```

```
print(paste("Lasso R-squared:", round(r2_lasso, 4)))
```

```
## [1] "Lasso R-squared: 0.961"
```

8.3 - Random Forest

```
# Random Forest  
# Why? Suitable for capturing nonlinear relationships and interactions between predictors.  
# Random Forest is also robust to overfitting and can rank feature importance.  
  
library(caret)
```

```

## Warning: package 'caret' was built under R version 4.3.3
## Loading required package: lattice
##
## Attaching package: 'lattice'
## The following object is masked from 'package:faraway':
##
##      melanoma
##
## Attaching package: 'caret'
## The following objects are masked from 'package:Metrics':
##
##      precision, recall
library(Metrics)
library(randomForest)

## Warning: package 'randomForest' was built under R version 4.3.3
## randomForest 4.7-1.2
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
## The following object is masked from 'package:dplyr':
##
##      combine
## The following object is masked from 'package:ggplot2':
##
##      margin
set.seed(139)
train_clean <- na.omit(train)
test_clean <- na.omit(test)

# Define the response variable
y_train <- train_clean$Deep..SWS..duration..min. / train_clean$Asleep.duration..min.
y_test <- test_clean$Deep..SWS..duration..min. / test_clean$Asleep.duration..min.

x_train <- train_clean %>% select(-Deep..SWS..duration..min.)
x_test <- test_clean %>% select(-Deep..SWS..duration..min.)

# Combine predictors and response for training
train_data <- cbind(x_train, y_train)

# Define cross-validation settings
set.seed(139)
cv_control <- trainControl(method = "cv", number = 5) # 5-fold cross-validation

# Train Random Forest model with cross-validation
rf_cv_model <- train(
  x = x_train,

```

```
y = y_train,
method = "rf",
trControl = cv_control,
tuneGrid = expand.grid(mtry = seq(2, sqrt(ncol(x_train)), length.out = 5)), # Adjust mtry values
ntree = 1000
)
```

```
## Warning: Setting row names on a tibble is deprecated.
```

```
## Warning: Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.  
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.
```

```
## Setting row names on a tibble is deprecated.  
## Setting row names on a tibble is deprecated.
```

```
# Print the best hyperparameters
```

```
print(rf_cv_model$bestTune)
```

```
##      mtry
```

```
## 4 3.769174
```

```
# Predict on test set using the best model
```

```
rf_cv_preds <- predict(rf_cv_model, x_test)
```

Evaluate the model: RMSE and R-squared

```
rmse_rf_cv <- rmse(y_test, rf_cv_preds)
```

```
rss_rf_cv <- sum((y_test - rf_cv_preds)^2)
```

```
tss_rf_cv <- sum((y_test - mean(y_test))^2)
```

```
r2_rf_cv <- 1 - (rss_rf_cv / tss_rf_cv)
```

```
print(paste("Cross-Validated Random Forest RMSE:", round(rmse_rf_cv, 4)))
```

```
## [1] "Cross-Validated Random Forest RMSE: 0.0373"
```

```
print(paste("Cross-Validated Random Forest R-squared:", round(r2_rf_cv, 4)))
```

```
## [1] "Cross-Validated Random Forest R-squared: 0.2516"
```